

AI Agent for Lunar Lander

Prabhjot Singh Rai (prabhjot) Abhishek Bharani (abharani)
Amey Naik (ameynaik)

December 13, 2018

1 Task Definition

The task accomplished by this project is to build an AI agent which would land with constant high rewards on a landing pad, defined by Box2D Lunar Lander available on OpenAI gym. This is accomplished by Reinforcement Learning, particularly by applying different Q-learning techniques. This project has explored Full DQN, Double DQN and Dueling Network Architecture (Dueling DQN), their performances in "solving" the game. We have considered a game to be solved when the agent starts getting average reward of 200 over 100 consecutive episodes. Moreover, performances have also been compared with baselines and oracle.

2 Infrastructure

We have used OpenAI library as our infrastructure. Although some insights are provided in Box2D Lunar Lander on OpenAI website, but thorough exploration of actions, state space, environment etc. were done before starting to solve the problem. Following is the description:

2.1 Actions

In this game, four discrete actions are available to the playing agent at any time frame:

- (a) Do nothing
- (b) Fire left orientation engine (rotates the lunar lander clockwise)
- (c) Fire main engine (provides upward thrust)
- (d) Fire right orientation engine (rotates the lunar lander anti-clockwise)

The agent can choose only one action among the given actions at a given time frame.

2.2 Terrain

The terrain is a combination of 10 points, and the helipad(landing zone) is fixed between 5th and 6th points towards the center. The values of the height of the landing zone(5th and 6th points on the terrain) are viewport height divided by 4, and the rest of the points are randomly sampled between 0 to $H/2$ using *numpy* random and smoothened (averaging 3 continuous points).

2.3 Initial State

The state is 8 dimensional values of different parameters of the lunar lander at any given time. The starting state is randomly initialized (the lunar lander takes a step in the world through the "idle" action) with certain bounds based on the environment.

Inspecting code of open_ai lunar lander, we see that the initial states are defined by simulating the environment in one frame (calling box world simulation using time step as $1/FPS$). Using this simulation, elaborating the initial state as follows:

1. Position X (Initial Position X: final x which changed from half of viewport width to a value after taking "idle" action before the simulation)
2. Position Y (Initial Position Y: final y which changed from from the viewport height to a value after taking "idle" action before simulation)
3. Velocity X (Initial Velocity X: final velocity x changed from 0 to a value after taking "idle" action before simulation)
4. Velocity Y (Initial Velocity Y: final velocity y which changed from 0 to a value after taking "idle" action before simulation)
5. Current lander angle (Initial lander angle: final lander angle after simulation on "idle" action from 0 degrees)
6. Angular velocity (Initial angular velocity: final angular velocity after simulation on "idle" action from 0 angular velocity)
7. Left leg contacted the surface (Initial value: False, since there's is no probability that the lunar lander's leg will touch the moon surface when at the top)
8. Right leg contacted the surface (Initial value: False)

2.4 End State

The episode ends in the following scenarios:

1. When the lunar lander goes outside of the viewport bounds, the game is over with -100 is negative reward.
2. When the lunar lander touches the ground with a high velocity
3. When the lunar lander touches the ground with body part except the legs
4. When the lunar lander stabilises on the moon's surface (change in shape of lunar lander is constantly 0 for a number of frames)

2.5 Rewards and Transitions

Before defining the rewards, let's define the shape of the lunar lander which decides the rewards. The shape of the lunar lander is a function of position coordinates (x, y) , linear velocities (v_x, v_y) , lander angle θ and contact of both the lander legs. We are interested in finding the change of shape at every step for the lunar lander to calculate the rewards for each given action. Shape change is given by subtraction of previous shape and current shape. Formally, shaping at time frame t :

$$\begin{aligned}
\text{shaping}_t &= -100 * (x^2 + y^2) \\
&\quad - 100 * (v_x^2 + v_y^2) \\
&\quad - 100 * \text{abs}(\theta) + 10 * (\text{Left leg contacted}) + 10 * (\text{Right leg contacted}) \\
\text{shape change} &= \text{shaping}_t - \text{shaping}_{t-1}
\end{aligned}$$

The rewards are defined as follows:

1. If the lunar lander crashes, or goes out of the bounds: -100
2. If the lunar lander is not awake anymore (stabilises at 0 shape change): $+100$
3. Doing nothing: shape change
4. Firing the engine: shape change - 0.3
5. Rotating: shape change - 0.03

The total reward will automatically be a sum of all the rewards at each time frame, and if the lander touches the ground with it's legs, will add those rewards to the total rewards earned during an episode. Transition probabilities are unknown, we get next states by simulating the lunar lander in the box environment given the current state and action taken.

Note: The transition probabilities and rewards are unknown to our agent, which it will try to figure out through exploration and incorporate in learning.

3 Models and different Approaches

3.1 DQN Introduction

Model-free based Deep Q Network algorithm was chosen specifically for the state size and complexity. DQN builds off of Q-learning algorithms by using a Deep Neural Network (DNN) for approximating the state-action value function, $Q(s, a)$. Function $Q(s, a)$ is defined such that for given state s and action a it returns an estimate of a total reward we would achieve starting at this state, taking the action and then following some policy. Lets call the Q function for the optimal policies Q_{opt} . Q_{opt} with discounting can be written as

$$Q_{opt}(s, a) = r_0 + \gamma r_1 + \gamma^2 r_2 + \dots \quad (1)$$

Here, r stands for rewards. γ is called a discount factor and when set it to $\gamma < 1$, it makes sure that the sum in the formula is finite. The γ controls how much the function Q in state s depends on the future and so it can be thought of as how much ahead the agent sees.

The above equation can be rewritten in a recursive form.

$$Q_{opt}(s, a) = r_0 + \gamma \max_a Q_{opt}(s', a) \quad (2)$$

This equation is proven to converge to the desired Q_{opt} , with finite number of states and each of the state-action pair is presented repeatedly. However, the Lunar Lander state space is real and continuous. We cannot store infinite number of values for every possible state. Instead, we approximate the Q function with a neural network. This network will take a state as an input and produce an estimation of the Q function for each action. This network with multiple layers is called Deep Q-network (DQN).

Experience Replay

During each simulation step, the agent perform an action a in state s , receives immediate reward r and come to a new state s' .

There are two problems with online learning -

1. The samples arrive in order they are experienced and as such are highly correlated. This might cause overfitting.
2. Throwing away each sample immediately after we use it. This means we are not using our experience effectively.

The key idea of 'experience replay' is that we store these (s, a, r, s') transitions in a memory and during each learning step, sample a random batch and perform a gradient descend on it. This way we solve both issues. After reaching the finite allotted memory capacity, the oldest sample is discarded. Exploration - **is this required?**

3.2 Full DQN

Full DQN - Separate Target network \tilde{Q} provides stable values and allows the algorithm to converge to the specified target:

$$Q(s, a) \rightarrow r + \gamma \max_a \tilde{Q}(s', a) \quad (3)$$

Theory: - In DQN algorithm we set targets for gradient descend as:

$$Q_{opt}(s, a) = r_0 + \gamma \max_a Q_{opt}(s', a) \quad (4)$$

We see that the target depends on the current network. A neural network works as a whole, and so each update to a point in the Q function also influences whole area around that point. And the points of $Q(s, a)$ and $\tilde{Q}(s, a)$ are very close together, because each sample describes a transition from s to s' . This leads to a problem that with each update, the target is likely to shift. This can lead to instabilities, oscillations or divergence.

To overcome this problem, researches proposed to use a separate target network for setting the targets. This network is a mere copy of the previous network, but frozen in time. It provides stable \tilde{Q} values and allows the algorithm to converge to the specified target:

$$Q(s, a) \rightarrow r + \gamma \max_a \tilde{Q}(s', a) \quad (5)$$

After several steps, the target network is updated, just by copying the weights from the current network. To be effective, the interval between updates has to be large enough to leave enough time for the original network to converge. **mention how often we are updating the model**
how often we are updating the model

3.3 Double DQN

Double DQN - Because of the max in the above formula, the action with the highest positive error is selected and this value is subsequently propagated further to other states. This leads to positive bias value overestimation.

$$Q(s, a) \rightarrow r + \gamma \max_a \tilde{Q}(s', \operatorname{argmax}_a Q(s', a)) \quad (6)$$

Theory: One problem in the DQN algorithm is that the agent tends to overestimate the Q function value, due to the max in the formula used to set targets. Because of the max in the formula, the action with the highest positive/negative error could be selected and this value might subsequently propagate further to other states. This leads to positive bias value overestimation. This severe impact on stability of the learning algorithm.

In this new algorithm, two Q functions Q_1 and Q_2 are independently learned. One function is then used

to determine the maximizing action and second to estimate its value. Either Q_1 or Q_2 is updated randomly with a formula:

$$Q_1(s, a) \rightarrow r + \gamma Q_2(s', \operatorname{argmax}_a Q_1(s', a)) \quad (7)$$

or

$$Q_2(s, a) \rightarrow r + \gamma Q_1(s', \operatorname{argmax}_a Q_2(s', a)) \quad (8)$$

It was proven that by decoupling the maximizing action from its value in this way, one can indeed eliminate the maximization bias.

When thinking about implementation into the DQN algorithm, we can leverage the fact that we already have two different networks giving us two different estimates Q and \tilde{Q} (target network). Although not really independent, it allows us to change our algorithm in a really simple way.

The original target formula would change to:

$$Q(s, a) \rightarrow r + \gamma \tilde{Q}(s', \operatorname{argmax}_a Q(s', a)) \quad (9)$$

how DDQN helped our lunar Lander game. were the values more stable?

3.4 Dueling layer DQN

- Dueling DQN - separate the estimators one that estimates the state value $V(s)$ one that estimates the advantage for each action $A(s, a)$

$$Q(s, a) \rightarrow A(s, a) + V(s) \quad (10)$$

$$Q(s, a; \theta, \alpha, \beta) \rightarrow V(s; \theta, \beta) + (A(s, a; \theta, \alpha) - \frac{1}{A} \sum_{a'} A(s, a'; \theta, \alpha)) \quad (11)$$

4 Experiments and Evaluation

In this section, we will present experimental results for all three different models presented in previous section as well as the visualization of training process.

4.1 Baselines

The first baseline is purely a random approach. The Agent was taking random actions and this is just to make sure our agent outperforms a random one. The second baseline is a simple linear classifier.

4.2 Training

In the most proposed models an initial learning rate of 10^{-4} was used to initialize training. Adam Optimizer is used for full duration of number of episodes. All models can be trained within 30-40 mins on CPU as input to the neural network is a state vector. All the models were trained for 800 episodes with batch size of 32 or 64. The number of episodes was chosen based on convergence of the loss, while batch size was chosen to be relatively small to get the benefits of stochasticity. We tried different epsilon decay rates to get the best results.

4.3 Hyperparameter Tuning

-We tested with different learning rate of 0.01, 0.002, 0.005 and 0.001. We noticed our model gave best result with 0.001

-We also tried different epsilon decay and got the best results at 0.995

-We tried various combinations of batchsize(32, 64 128). Of all combinations, we saw best score with batch size of 64 across all models.

Model	Avg. Score	Number of episodes to reach 200 score
Baseline	-200	never
Linear Baseline	-150	never
DQN	123	525
Double DQN	220	500
Dueling DQN	225	435

Table 1: Comparison of Results for different model implementations




Figure 1: Rewards for different approaches on tensor board




Figure 2: Different Sets of hyperparamerts for DQN Model

4.4 Analysis

Error analysis: After doing hyper-parameter tuning, our Agent was able to achive average score of more than 200 quickly (at 435 episodes). However for some episodes after 435 the rewards were not consistently more than 200 for 100 iterations and we observe the variation is more in DQN and Double DQN as compared to Dueling DQN. We got the best performance on Set 4 with Dueling DQN model.

4.5 Model Evaluation

All aforementioned models are evaluated for 100 episodes after 450 episodes.

Hyper-parameter	Set 1	Set 2	Set 3	Set 4
gamma	0.99	0.99	0.99	0.99
Epsilon(max,min,decay)	(1,0,.998)	(1,.01,.995)	(1,.01,.998)	(1,.01,.998)
Learning Rate	0.001	0.0001	0.0001	0.0001
DNN layers	[32,32]	[128,32]	[128,64]	[128,64]
Loss Function	MSE	MSE	MSE	MSE
Batch Size	32	32	64	64
Replay Memory Size	2^{16}	2^{16}	2^{16}	2^{16}

Table 2: Different set of hyper-parameters were tried




Figure 3: Evaluating Performance of different DQN Networks

5 Conclusion

6 Codalab Link

<https://worksheets.codalab.org/worksheets/0x3f15e8eba6af45828a26df3bc0e1f490/>
command to run

```
cl run :brain.py :hyperparameters.py :run.py :agent.py :environment.py :memory.py 'python3 run.py' -n
sort-run --request-docker-image prabhjotrai/openai-gym:v1
```